

# Abstract Semantic Differencing for Numerical Programs

## Abstract

We address the problem of computing semantic differences between a program and a patched version of the program. Our goal is to obtain a precise characterization of the difference between program versions, or establish their equivalence when no difference exists.

We focus on computing semantic differences in numerical programs where the values of variables have no a-priori bounds, and use abstract interpretation to compute an over-approximation of program differences. Computing differences and establishing equivalence under abstraction requires abstracting relationships between variables in the original program and its patched version. Towards that end, we first construct a *union program* in which these relationships can be identified, and then use a *correlating abstract domain* to compute a sound approximation of these relationships. To establish equivalence between correlated variables and precisely capture differences, our domain has to represent non-convex information. To balance precision and cost of this representation, our domain may over-approximate numerical information as long as equivalence between correlated variables is preserved.

We have implemented our approach in a tool built on the LLVM compiler infrastructure and the APRON numerical abstract domain library, and applied it to a number of challenging real-world examples, including programs from the GNU core utilities, [Mozilla Firefox and the Linux Kernel](#). We evaluate over 50 patches and show that for these programs, the tool often manages to establish equivalence, reports useful approximation of semantic differences when differences exist, and reports only a few false differences.

## 1. Introduction

When applying a patch to a procedure, the programmer has very limited means for acquiring a description of the change the patch made to the procedure behavior.

Given a program  $P$  and a patched version of the program  $P'$  our goal is to determine the difference between  $\llbracket P \rrbracket$  and  $\llbracket P' \rrbracket$ .

Our goal is to describe semantic difference between two programs (program versions)  $P, P'$  at certain program points i.e. given two programs and two locations in the programs, we want to produce all program states that are possible at the given location in one program but do not exist in the other program and vice versa.

**Existing Techniques** Existing techniques for proving patch equivalence will only supply the programmer with a binary answer [4] as to the (input-output) equivalence a program but no description of the difference is supplied. Further work [5] allows refining the equivalence proof by providing a set of constraints under which equivalence is desired but requires the programmer to manually deduce these. Other techniques for describing difference [?] which rely on symbolic execution supply unsound results as they are limited by loops and essentially cover a subset of program behavior. We present a novel approach which allows for a sound description of difference for programs with loops. Our technique employs methods of abstract interpretation for over-approximating the difference in behaviors, by focusing the abstraction on the *relationships*

between program behaviors i.e. between variables values (data) and conditionals (path) in the two versions.

In contrast to existing techniques, our approach allows checking for equivalence in every point of execution, and for every variable, while previous approaches focus only on input-output equivalence. This enables detection of key differences that impact the correctness of the patch: if the changed behavior includes a bug manifested by a local variable (for instance: array index out of bounds), we will detect and describe it while previous work only detected it when propagated to the output and equivalence may have been reported although a bug was introduced. This also provides a challenge as we need to carefully choose the program locations where we check for difference otherwise we will spuriously detect difference.

**Correlating Program** Abstracting relationships allows us to maintain focus on difference while omitting (whenever necessary for scalability) parts of the behavior that does not entail difference. In order to monitor these relationships we created a *correlating program* which captures the behavior of both the original program and its patched version. Instead of designing a correlating semantics that is capable of co-executing two programs, we chose to automatically construct the correlating program such that we can benefit from the use of standard analysis frameworks for analyzing the resulting program. Another advantage of this new construct, is that you may apply other methods for equivalence checking directly on it [8] as the correlation allows for a much more fine-grained equivalence checking (between local variables and not only output).

**Correlating Abstractions** Our abstraction holds data of both sets of variables, joined together and is initialized to hold equality over all matched variables. This means we can reflect relationships without necessarily knowing the actual value of a variables (we can know that  $x_{old} = x_{new}$  even though actual values are unknown). We ran out analysis over the correlating program while updated the domain to reflect program behavior.

To establish equivalence between correlated variables and precisely capture differences, our domain has to maintain correlating information even when other information is abstracted away.

Since some updates may result in non-convex information (e.g. taking a condition of the form  $x \neq 0$  into account), our domain has to represent non-convex information, at least temporarily. We address this by working with a powerset domain of a convex representation. To avoid exponential blowup, our join operator may over-approximate numerical information as long as equivalence between correlated variables is preserved.

In some cases, it would have been sufficient to use alternative domains that are capable of representing richer information, such as interval polyhedra [3], or other numerical domains that can represent non-convex information (e.g., [?]). The recent donut domain [?] may be of particular interest for this purpose. However, the general principle of having to preserve correlating information even when information about the values is abstracted away, holds in all of these cases.

```

int sign(int x) {
  int sgn;
  if (x < 0)
    sgn = -1
  else
    sgn = 1
  return sgn
}

int sign'(int x) {
  int sgn;
  if (x < 0)
    sgn = -1
  else if (x==0)
    sgn = 0
  else
    sgn = 1
  return sgn
}

```

Figure 1: Two simple implementations of the *sign* operation.

In this paper, we present a technique based on abstract interpretation, that is able to compute an over-approximation of the difference between numerical programs or establish their equivalence when no difference exists. The approach is based on two key ideas: (i) create a *union program* that captures the behavior of both the original program and its patched version; (ii) analyze the union program with a *correlating domain* that captures relationships between values of variables in the original program and values of variables in the patched version.

The idea of a correlating program is similar to that of self-composition [2, 10], but the way in which statements in the union program are combined is carefully designed to keep the steps of the two programs close to each other. Rather than having the patched program sequentially composed after the original program, our union program interleaves the two versions. Analysis of the union program can then recover equivalence between values of correlated variables even when equivalence is *temporarily* violated by an update in one version, as the corresponding update in the other version follows shortly thereafter.

**check whether Aiken paper SAS'05 does some sort of interleaving as well**  
**also need to say that there is the problem of choosing differencing points**

## 1.1 Main Contributions

The main contributions of this paper are as follows:

- we phrase the problem of semantic differential analysis as an analysis of a union program — a single program that represents an original program and its patched version.
- we present an approach for analyzing differences over the union program using a correlating abstract domain. Our approach is sound — if there is a difference at a differentiation point, we cannot miss it. However, since we over-approximate differences, our approach may report false differences due to approximation.
- We have implemented our approach in a tool based on the LLVM compiler infrastructure and the APRON numerical abstract domain library, and evaluated it using over 50 patches from open-source software including GNU core utilities, Mozilla Firefox, and the Linux Kernel. Our evaluation shows that the tool often manages to establish equivalence, reports useful approximation of semantic differences when differences exists, and reports only a few false differences.

**mention classical work like [? ? ]**

## 2. Overview

In this section, we informally describe our approach with a few simple example programs.

```

int sign(int x) {
  int x' = x;
  Guard G1 = (x < 0);
  Guard G1' = (x' < 0);
  int sgn;
  int sgn';
  if (G1) sgn = -1;
  if (G1') sgn' = -1;
  Guard G2' = (x' == 0);
  if (!G1') if (G2') sgn' = 0;
  if (!G1') if (!G2') sgn' = 1;
  if (!G1) sgn = 1;
}

```

Figure 2: Correlating program  $\text{sign} \bowtie \text{sign}'$ .

### 2.1 Motivating Example

Consider the simple example program of Fig. 1, inspired by an example from [? ]. For this example, we would like to establish that *sign* and *sign'* only differ in the case where  $x = 0$ .

As a first naive attempt one could try to analyze each version of the program separately and compare the (abstract) results. However, this is clearly unsound, as equivalence under abstraction does not entail concrete equivalence. For example, using an interval analysis [? ] would yield that in both programs the value of *sgn* ranges in the same interval  $[-1, 1]$ , missing the fact that *sign* never returns the value 0.

**Establishing equivalence under abstraction** To establish equivalence under abstraction, we need to abstract relationships between the values of variables in *sign* and *sign'*. Specifically, we need to track the relationship between the values of *sgn* in both versions and see whether we can establish their equivalence.

We reduce the problem of analysis across two programs to the problem of analyzing a single program by constructing a *correlating program*. The correlating program represents the behaviors of both programs, and allows us to reason about relationships between variables in both.

Fig. 2 shows the correlating program for the programs of Fig. 1. Using the correlating program, we can directly track the relationship between *sgn* in *sign* and its corresponding variable *sgn'* in *sign'*. Unfortunately, any domain with convex constraints will still fail to capture the precise relationship between *sgn* and *sgn'*. For example, using the polyhedra abstract domain [? ], the relationship between *sgn* and *sgn'* in the correlating program would be **(constraint – here)**.

An obvious, but prohibitively expensive, solution to the problem is to use disjunctive completion [? ], moving to a powerset domain in which every abstract state is a set of convex objects (e.g., set of polyhedra). However, using such domain would significantly limit the applicability of the approach.

The desirable solution is a partially disjunctive domain, in which only certain disjunctions are kept separate during the analysis, while others are merged. The challenge in our setting is in keeping the partition fine enough such that equivalence could be preserved, without reaching exponential blowup.

**How to partition** We base our choice of explicit disjunctions to be kept based on their effect on variable equivalence.

### Widening

## 3. Preliminaries

We use the following standard concrete semantics definitions for a program:

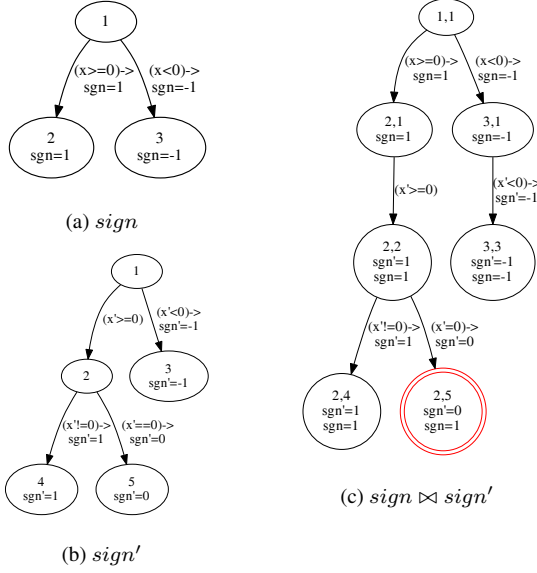


Figure 3: Transition systems for the programs of Fig. 1 and Fig. 2.

**Program Location / Label** A program location  $loc \in Loc$ , also referred to as label denoted  $lab$ , is a unique identifier for a certain location in a program corresponding to the value of the program counter at a certain point in the execution of the program. We also define two special labels for the start and exit locations of the program as *begin* and *fin* respectively.

**Concrete State** Given a set of variables  $Var$ , a set of possible values for these variables  $Val$  and the set of locations  $Loc$ , a *concrete program state* is a tuple  $\sigma \triangleq \langle loc, values \rangle \in \Sigma$  mapping the set of program variables to their concrete value at a certain program location  $loc$  i.e.  $values : Var \rightarrow Val$ . The set of all possible states of a program  $P$  is denoted  $\Sigma_P$ .

**Program** We describe an imperative program  $P$ , as a tuple  $(Val, Var, \rightarrow, \Sigma_0)$  where  $\rightarrow : \Sigma_P \times \Sigma_P$  is a transition system which given a concrete program state returns the following state in the program and  $\Sigma_0$  is a set of initial states of the program. Our formal semantics need to deal with errors states therefore we ignore crash states of the programs, as well as inter-procedural programs since our work deals with function calls by inlining and we exclude recursion for now.

**Concrete Trace** A program trace  $\pi \in \Sigma_P^*$ , is a sequence of states  $\langle \sigma_0, \sigma_1, \dots \rangle$  describing a single execution of the program. Each of the states corresponds to a certain location in the program where the trace originated from. Every program can be described by the set of all possible traces for its run  $\llbracket P \rrbracket \subseteq \Sigma_P^*$ . We refer to these semantics as concrete state semantics. We also define the following standard operations on traces:

- $label : \Sigma_P \rightarrow Lab$  maps a state to the program label at which it appears.
- $last : \Sigma_P^* \rightarrow \Sigma_P$  returns the last state in a trace.
- $pre : \llbracket P \rrbracket \rightarrow 2^{\Sigma_P^*}$  for a trace  $\pi$  is the set of all prefixes of  $\pi$ .

## 4. Concrete Semantics

In this section, we define the notion of difference between concrete semantics of two programs based on a standard concrete semantics for a program.

### 4.1 Concrete State Differencing

Comparing two different programs  $P$  and  $P'$  under our concrete semantics means comparing *traces*. A trace is composed of concrete states thus we first need a way to compare states. We define a state  $\sigma$  as a mapping  $Var \rightarrow Val$ , and a difference between states would be one state mapping a different value to the same variable. But as variables may differ between programs, we require a correspondence between variables in  $P$  (i.e.  $Var$ ) and those in  $P'$  (i.e.  $Var'$ ) and the difference will be checked over matched variables.

**Variable Correspondence** A variable correspondence  $VC \in Var \times Var'$ , is a partial mapping between 2 sets of program variables. Any variable in  $Var$  may be matched with any in  $Var'$  and vice versa. Naturally, the comparison between states will occur between values mapped to corresponding variables, as described by  $VC$ .

Our analysis may receive a user defined correspondence or use a standard name-base matching:  $VC_{EQ} \triangleq \{(v, v') | v \in Var \wedge v' \in Var' \wedge name(v) = name(v')\}$  and vice versa. Our experience suggests that for most patches, the set of variables does not change over subsequent versions ( $Var = Var'$ ) except for perhaps the addition of new variables or removal of an old one. Therefore we concluded that for our purposes, matching variables by name is sufficient for finding a precise difference.

**Concrete State Delta** Given two concrete states  $\sigma \in \Sigma_P, \sigma' \in \Sigma_{P'}$  and a variable correspondence  $VC$ , we define the concrete state delta  $\Delta_S(\sigma, \sigma') : \Sigma_P \times \Sigma_{P'} \rightarrow 2^{Var \times Val}$  as the part of the state  $\sigma$  where corresponding variables do not agree on values (with respect to  $\sigma'$ ). Formally:  $\Delta_S(\sigma, \sigma') \triangleq \{(var, val) | (var, var') \in VC \wedge \sigma(var) = val \neq \sigma'(var')\}$ . In case there is no observable difference in state we get that  $\Delta_S(\sigma, \sigma') = \emptyset$ . As state delta is directional, we notice that unless it is empty,  $\Delta_S$  is not symmetric. In fact, the direction in which  $\Delta_S$  is used has meaning in the context of a program  $P$  and a patched version of it  $P'$ . We define  $\Delta_S^- = \Delta_S(\sigma, \sigma')$  which means the values of the state that was "removed" in  $P'$  and  $\Delta_S^+ = \Delta_S(\sigma', \sigma)$  which stands for the values added in  $P'$ .

**Example** For instance, given two states  $\sigma = (x \mapsto 1, y \mapsto 2, z \mapsto 3)$  and  $\sigma' = (x' \mapsto 0, y' \mapsto 2, w' \mapsto 4)$  and using  $VC_{EQ}$  then  $\Delta_S^- = (x \mapsto 1)$  since  $x$  and  $x'$  match and do not agree on value,  $y$  and  $y'$  agree (thus are not in delta) and  $z'$  is not in  $VC_{EQ}$ . Similarly,  $\Delta_S^+ = (x' \mapsto 0)$ .

We defined a notion for difference between states but this is insufficient to describe difference between program traces. Program equivalence is defined as behavior, or output, equivalence given **the same input**. Therefore we are only interested in comparing traces that originate from the same input (for the input variables correlated by  $VC$ ). Every mention of trace differentiation will assume the traces agree on input. The way to differentiate traces is by differentiating their states, but which states? this is not a trivial question since traces can vary in length and order of states. We need a mapping for choosing the states to be differentiated within the two traces. We define it as following:

**Trace Diff Points** Given two traces  $\pi \in \llbracket P \rrbracket$  and  $\pi' \in \llbracket P' \rrbracket$  that originate from the same input, we define a trace index correspondence relation named trace diff points denoted  $DP_\pi$  as a matching of indexes specifying states where concrete state delta should be computed. Formally:  $DP_\pi \subseteq \{(i, i') | i \in 0..|\pi|, i' \in 0..|\pi'|\}$ . The

```

void foo(unsigned x) {      void foo'(unsigned x) {
    unsigned i = 0;          unsigned i = 0;
    lab:guard g = (i >= x);  lab:guard g = (i >= 2*x);
    if (g) return;          if (g) return;
    ...                     ...
    i++;                    i++;
    goto lab;               goto lab;
}                          }

```

Figure 4:  $P, P'$  differentiation candidates

question of supplying this matching, in a way that results in meaningful delta, is not a trivial one, we delay this discussion until we define the trace delta.

Now that we have a way of matching states to be compared between two traces, we define the notion of trace differentiation:

**Trace Delta** Given two traces  $\pi \in \llbracket P \rrbracket$  and  $\pi' \in \llbracket P' \rrbracket$  that originate from the same input, and a state correspondence  $DP_\pi$  we define the trace delta  $\Delta_T(\pi, \pi') : \llbracket P \rrbracket \times \llbracket P' \rrbracket \rightarrow (Loc \rightarrow 2^{Var \times Val})$  as state differentiations between all corresponding states in  $\pi$  and  $\pi'$ . Formally, for every  $(i, i') \in DP_\pi$  such that  $\Delta_S(\sigma_i, \sigma_{i'}) \neq \emptyset$ ,  $\Delta_T$  will contain the mapping  $i \mapsto \Delta_S(\sigma_i, \sigma_{i'})$ , thus the result will map certain states in  $\pi$  to their state delta with the corresponding state<sup>+</sup> in  $\pi'$  (deemed interesting by  $DP_\pi$ ). We define  $\Delta_T^+$  and  $\Delta_T^-$  in a similar way.

One possible choice for  $DP_\pi$  would be the endpoints of the two traces  $\{(fin, fin')\}$  (assuming they are finite) meaning differentiating the final states of the executions or formally:  $\Delta_{fin}^- \triangleq \Delta_{fin}(\pi, \pi') = \{fin \mapsto \Delta_S(\sigma_{fin}, \sigma'_{fin'})\}$  (we will also be interested in  $\Delta_{fin}^+$ ). The final state delta may not always be sufficient for truly describing the difference between said traces as according to our definition of difference, we would be interested in checking difference at intermediate locations in the program (that emit output or check assertions).  $\Delta_{fin}$  will only compare final state values and could miss what happened during the execution. This can be overcome by instrumenting the semantics such that the state will contain all "temporary" values for a variable along with the trace index (program location) is not sufficient here as a trace can loop over a certain location) where the values existed by, for instance, adding temporary variables, allowing for a complete differentiation at the end point. This may substantially complicate the selection of  $VC$  as it will require a matching between all of these temporary, indexed, variables. Such a correspondence may be extremely hard to produce. Also the number of variables here can range up to the length of the trace.

Defining the  $DP$  over any two traces is a daunting task since traces of separate (although similar) programs can vastly differ. If we take a look at two versions of a program depicted in Fig. ?? and the following traces generated from the input  $x = 2$ :  $\pi = \langle (x \mapsto 2, i \mapsto 0, g \mapsto 0), (x \mapsto 2, i \mapsto 1, g \mapsto 0), (x \mapsto 2, i \mapsto 2, g \mapsto 1) \rangle$  and  $\pi' = \langle (x \mapsto 2, i \mapsto 0, g \mapsto 0), (x \mapsto 2, i \mapsto 1, g \mapsto 0), (x \mapsto 2, i \mapsto 2, g \mapsto 0), (x \mapsto 2, i \mapsto 3, g \mapsto 0), (x \mapsto 2, i \mapsto 4, g \mapsto 1) \rangle$ , we see that even in this simple program, finding a correlation based on traces alone is hard. However, one can get the sense that using program location as a means of correlation, one can produce a meaningful result that describes how the values of  $i$  range differently in the new version  $P$ . For example, if we look at all the possible values for  $i$  in label  $lab$  and differentiate them (as a set) from the values in the patched version (in the same location), we get a meaningful result that  $i$  in the patched version can range from  $x+1$  up to  $2x$ . We will discuss differentiation of sets of states later on as we describe the collecting semantics.

## 4.2 Differencing at Program Labels

**Trace Delta using Program Label** Given two traces  $\pi, \pi'$  and two program labels  $l, l'$  we define a trace delta based on all trace locations (states) that are labeled  $l, l'$ . First we define  $\pi_l$  as a subsequence of  $\pi$  where only states that are labeled  $l$  were chosen ( $\pi_{l'}$  is defined similarly). Next, we denote  $\Delta_L(\pi_l, \pi_{l'})$  as a means for comparing these sequences. As  $\pi_l, \pi_{l'}$  may vary in length and order, we cannot simply define it as applying  $\Delta_S$  on each pair of states in  $(\pi_l, \pi_{l'})$  by order. In fact,  $\Delta_L$  can be defined in different way to reflect different concepts of difference, for instance, it can be defined as the differentiating the last states of  $\pi_l$  and  $\pi_{l'}$  (assuming they are both finite) to reflect we are only interested in the final values in that location. We chose to define  $\Delta_L$  as the difference between the set of states which appear in  $\pi_l$  against the set of those in  $\pi_{l'}$ . Formally:  $\Delta_{L_{set}}(\pi_l, \pi_{l'}) \triangleq \{\sigma \in \text{ran}(\pi_l) \mid \neg \exists \sigma' \in \text{ran}(\pi_{l'}) \text{ s.t. } \Delta_S(\sigma, \sigma') = \emptyset\}$ . For example consider Fig. ??, for  $\pi, \pi'$  that originate from  $x = 2$  then  $\Delta_{L_{set}}(\pi_{lab}, \pi'_{lab'}) = \emptyset$  and  $\Delta_{L_{set}}(\pi'_{lab'}, \pi_{lab}) = \{(i' \mapsto 3), (i' \mapsto 4)\}$ . We see that this notion of  $\Delta$  indeed captures a useful description of difference.

The problem of choosing  $DP$  is now reduced to the matching of labels as the trace indexing correspondence  $DP_\pi$  defined in Definition 4.1 is induced by the definition over labels. as we need to differentiate sets of states belonging to a certain program label. We require a correspondence of labels and therefore we define the label diff points correspondence.

**Label Diff Points** Given two programs  $P, P'$  and their sets of program labels  $Lab, Lab'$ , we define a label correspondence relation named label diff points denoted  $DP_{Lab}$  as a matching of labels between programs. Formally:  $DP_{Lab} \subseteq \{(l, l') \mid l \in Lab, l' \in Lab'\}$ . From this point on any mention of the diff-points correspondence  $DP$  will refer to label diff-points  $DP_{Lab}$ . We address the question of selection of  $DP_{Lab}$  in ??.

Now, we will move past the concrete semantics towards a *collecting semantics* as a step towards abstraction. This is required as it is unfeasible to describe difference based on traces. However, we need to adjust our concrete semantics before we can correctly define this as the collecting semantics based on individual traces **will not allow us to correlate traces that originate from the same input**. This is the first formal indication of how a separate abstraction, that considers each of the programs by itself, cannot succeed.

## 4.3 Concrete Correlating Semantics

In this section we shortly define, based on previous definitions, the correlating state and trace which bind the executions of both programs,  $P$  and  $P'$ , together and we define the notion of delta there. Essentially, these will be states and traces of the product program  $P \times P'$  but **only traces that originate from the same input are considered**. This allows us to then define the correlating abstract semantics which is key for successful differencing.

**Correlating Concrete State** A correlating concrete state  $\sigma_\times$  is a pair of concrete states  $\langle \sigma, \sigma' \rangle$  belonging to a pair of programs  $P$  and  $P'$ . The set of all possible correlating states is denoted  $\Sigma_{P \times P'}$ .

**Correlating Concrete Trace** A correlating trace  $\pi_\times$ , is a sequence of correlating states  $\dots, \langle \sigma_i, \sigma'_j \rangle, \dots$  describing a dual execution of the two programs where at every point each of the program can perform a single step. The  $label_\times$ ,  $last_\times$  and  $pre_\times$  operations are defined similarly. We denote by  $\llbracket P \times P' \rrbracket$  the set of all traces in  $P \times P'$ . Again, we restrict to traces that originate from the same input i.e.  $\sigma_0 =_{VC} \sigma'_0$ .

We must remember however, that the number of traces to be compared is potentially unbounded which means that the delta we compute may be unbounded too. Therefore we must use an abstraction over the concrete semantics that will allow us to represent executions in a bounded way.

## 5. Abstract Correlating Semantics

In this section, we introduce our correlating abstract domain which allows bounded representation of product program state while maintaining equivalence between correlated variables. This comes at the cost of an acceptable loss of precision of other state information. We represent variable information using standard relational abstract domain. As our analysis is path sensitive, we allow for a set of abstract sub-states, each adhering to a certain path in the product program. This abstraction is similar to the trace partitioning domain as described in [9]. To assure we only consider correlated paths from the product programs (that agree on input), our abstraction will initially assume equality on all inputs. This power-set domain records precise state information but does not scale due to exponential blow-up of number of paths. To reduce state size, we define a special join operation that dynamically partitions the abstract state according to the set of equivalences maintained in each sub-state and joins all sub-states in the same partition together (using the sub-domain lossy join operation). This equivalence criteria allows separation of equivalence preserving paths thus achieving better precision. We start off by abstracting the correlating trace semantics in ??.

In the following, we assume an abstract relational domain  $(D^\sharp, \sqsubseteq_D)$  equipped with operations  $\sqcap_D$ ,  $\sqcup_D$  and  $\nabla_D$ , for representing sets of concrete states in  $\Sigma_{P \times P'}$ . We separate the set of program variables into original program variables denoted  $Var$  and the added guard variables denoted  $Guard$  that are used for storing conditional values alone. We assume the abstract values in  $D^\sharp$  are constraints over the variables and guards (we denote  $D^\sharp_{Guard}$  for abstraction of guards and  $D^\sharp_{Var}$  for abstracting original variables), and do not go into further details about the particular abstract domain as it is a parameter of the analysis. We also assume that the sub-domain  $D^\sharp$  allows for a sound over-approximation of the concrete semantics (given a sound interpretation of program operations).

**Correlating Abstract State** A correlating abstract program state  $\sigma^\sharp \in Lab_\times \rightarrow 2^{D^\sharp_{Guard} \times D^\sharp_{Var}}$ , is a set of pairs  $\langle ctx, data \rangle \in \Sigma^\sharp$  mapped to a product program label  $l_\times$ , where  $ctx \in D^\sharp_{Guard}$  is the execution context i.e. an abstraction of guards values via the relational numerical domain and  $data \in D^\sharp_{Var}$  is an abstraction of the variables (anything that is not a guard) also using the domain  $D^\sharp$ . We separate abstractions over guard variables added by the transformation to GCL and original program variables as there need not be any relationships between guard and regular variables.

**Correlating Abstract Semantics** We describe the abstract semantics transformers. Given an abstract state  $\sigma^\sharp = l_\times \mapsto S$ :

- $\llbracket v := e \rrbracket^\sharp = l_\times \mapsto \{ \langle ctx, \llbracket v := e \rrbracket^\sharp_{D^\sharp}(data) \rangle \mid \langle ctx, data \rangle \in S \}$
- $\llbracket g := e \rrbracket^\sharp = l_\times \mapsto \{ \langle \llbracket g := true \rrbracket^\sharp_{D^\sharp}(ctx), \llbracket e \rrbracket^\sharp_{D^\sharp}(data) \rangle \mid \langle ctx, data \rangle \in S \} \cup \{ \langle \llbracket g := false \rrbracket^\sharp_{D^\sharp}(ctx), \neg e \rrbracket^\sharp_{D^\sharp}(data) \rangle \mid \langle ctx, data \rangle \in S \}$
- $\llbracket \text{if}(g)\{s_0\}\text{else}\{s_1\} \rrbracket^\sharp = l_\times \mapsto \{ \langle \llbracket g := true \rrbracket^\sharp_{D^\sharp}(ctx), \llbracket s_0 \rrbracket^\sharp_{D^\sharp}(data) \rangle \mid \langle ctx, data \rangle \in S \} \cup \{ \langle \llbracket g := false \rrbracket^\sharp_{D^\sharp}(ctx), \llbracket s_1 \rrbracket^\sharp_{D^\sharp}(data) \rangle \mid \langle ctx, data \rangle \in S \}$
- $\llbracket \text{goto } lab \rrbracket^\sharp = \sigma^\sharp$

We assume that any program  $P$  can be transformed such that it contains only the aforementioned operations (specifically GCL format). We also assume that for  $\llbracket g := e \rrbracket^\sharp$  operations,  $e$  is a logical operation with binary value. Next, we define the abstraction function  $\alpha : 2^{\Sigma_{times}} \rightarrow 2^{D^\sharp \times D^\sharp}$  for a set of concrete traces  $T \subseteq \Sigma_{times}^*$ . As in our domain traces are abstracted together

```

int f(int x) {      int f'(int x) {
    return x;      return 2*x;
}                  }

```

caption: single path differentiation candidates

if they share the exact same path, we first define an operation  $path : \Sigma_{times}^* \rightarrow Lab^*$  which returns a sequence of labels for a trace's states i.e. what is the path taken by that trace. We also allow applying  $path$  on a set of traces to denote the set of paths resulting by applying the function of each of the traces. Finally:  $\alpha(T) \triangleq \{ \sqcup_{path(\pi)=p} \beta(last(\pi)) \mid p \in path(T) \}$  where  $\beta(\sigma^\sharp) = \langle \beta_{D^\sharp}(\sigma^\sharp|_{Guard}), \beta_{D^\sharp}(\sigma^\sharp|_{Var}) \rangle$  i.e. applying the the abstraction function of the abstract sub-domain  $\beta_{D^\sharp}$  on parts of the concrete state applying to  $Guards$  and  $Vars$  separately. Our abstraction partitions trace prefixes  $\pi$  by path and abstracts together the concrete states reached by the prefix -  $last(\pi)$ , using the sub-domain.

Every path in the product program will be represented by a single abstract of the sub-domain. As a result, all **traces prefixes** that follow the same path to  $l_\times$  will be abstracted into a single sub-state of the underlying domain. This abstraction fits semantics differencing well, as inputs that follow the same path display the same behavior and will usually either keep or break equivalence together, allowing us to separate them from other behaviors (it is possible for a path to display both behaviors as in Fig. ?? and we will discuss how we are able to manipulate the abstract state and separate equivalent behaviors from ones that offend equivalence). Another issue to be addressed is that fact that our state is still potentially unbound as there may be an infinite number of paths in the program (due to loops).

**Canonization** Performing analysis with the powerset domain does not scale as the number of paths in the correlated program may be exponential. We must allow for reduction of state  $\sigma^\sharp = l_\times \mapsto S$  with acceptable loss of precision. This reduction, or canonization as we call it, can be achieved by joining the abstract sub-states in  $S$  (using the standard precision losing join of the sub-domain) but to perform this we must first answer the following: (i) Which of the sub-states shall be joined together i.e. what is the *Canonization Strategy* and (ii) At which program locations should the canonization occur i.e. what is the *Canonization Point*. A trivial canonization strategy (we name *Join-All*) is simply reverting back to the sub-domain by applying the join on all sub-states which may result in unacceptable precision loss as exemplified in Fig. 2. However, by taking a closer look at the final state of the same example  $\sigma^\sharp(fin) = [ \langle (g1 = 1, g2' = 0, \equiv_{g1}), (sgn = 1, \equiv_{x,sgn}) \rangle, \langle (g1 = 0, g2' = 0, \equiv_{g1}), (x < 0, sgn = -1, \equiv_{x,sgn}) \rangle, \langle (g1 = 0, g2' = 1, \equiv_{g1}), (x = 0, sgn = 0, sgn' = 1, \equiv_x) \rangle ]$ , one may observe that were we to join the two sub-states that maintain equivalence on  $\{x, sgn, g1\}$ , it would result in an acceptable loss of precision (of losing the  $x$  related constraints). This led us to devise a canonization strategy (we name *Join-Equiv*) that partitions sub-states by the set of variables which they preserve equivalence for. This bounds the super-state size at  $2^{|V|}$ , where  $V$  is the set of correlating variables we wish to track. As mentioned, another key factor in preserving equivalence and maintaining precision is the program location at which the canonization occurs. The first possibility, which is somewhat symmetric to the first canonization strategy, is to canonicalize at every join point i.e. after every branch converges. We name this canonization point selection as *At-Join*. A quick look at Fig. 2's state after processing the first guarded instruction `if (G1) sgn = -1;` i.e.  $\sigma^\sharp = [ \langle (g1 = 0, \equiv_{g1}), (x \geq 0, \equiv_{x,sgn}) \rangle, \langle (g1 = 1, \equiv_{g1}), (x < 0, sgn' = -1, \equiv_x) \rangle ]$ , (we ignored  $g2'$  effect at this point for brevity) suggests that applying canonization strategy 1, will perform badly in some scenerios, specifically here as we will lose all data regarding  $sgn$  (remem-



ber, we do not assume anything regarding canonization strategy therefore the canonization by equivalence will not necessarily save us here). However if we could delay the canonization to a point where the two programs "converge" (after the following `if (G1') sgn' = -1; line;`), we will get a more precise temporary result which preserves equivalence. Therefore we define another canonization point we call *At-Converge* which are defined during the correlating program construction process described in Section 6.4 and are basically locations where the two (guarded) programs have syntactically converged (we found two lines that are syntactically equal, sans tagging). A final strategy we define uses our *differencing points* as canonization points and delay applying canonization until then. We name this canonization point selection as *At-Diff*. We remind that diff-points are product program locations where both programs conceptually converge as they are about to emit output i.e. both programs have finished "preparing" the next portion of output therefore, if equivalence exists, it must exist now, therefore this is a prime candidate for a canonization point. Our evaluation includes applying each of our strategies along with each of the canonization points. Intuitively, the results should range from least precise using the  $\langle \text{Join-All}, \text{At-Join} \rangle$  strategy and point to most precise in the  $\langle \text{Join-Equiv}, \text{At-Diff} \rangle$  scenario and this is indeed the case as we show in Section 7 (not taking into account the no-canonization scenario which is naturally most precise).

**Widening** In order for our analysis to handle loops we require a means for reaching a fixed point. As our analysis advances over a loop and state is transformed, it may keep changing and never converge unless we apply the widening operator to further over-approximate the looping state and arrive at a fixed point. We have the widening operator of our sub-domain at our disposal, but again we are faced with the question of how we apply this operator, i.e. which pairs of sub-states  $\langle \langle \text{ctx}, \text{data} \rangle \rangle$  from  $\sigma^\#$  should be widened with which  $\sigma'_j$  in  $\Sigma'$  or which *Widening Strategy* should we apply. A first viable strategy, similar to the first canonization strategy, is to perform an overall join operation on all  $\sigma_i$  and  $\sigma'_j$  to result in a single convex sub-state and then simply applying the widening to these sub-states using the sub-domain's  $\nabla$  operator. If we examine applying this strategy to  $P \bowtie P1$  from Fig. ??, we see that it may successfully arrive at a fixed point that also maintains equivalence. Now let us try applying the strategy to  $P \bowtie P2$  as seen in Fig. ?. First we mention that as P2 introduces a return statement under the  $len = 0$  condition, the example shows an extra  $r$  guard for representing a return (this exists in all GCL programs but we omitted it so far for brevity). While analyzing, once we pass line ??, our state is split to reflect the return effect  $\Sigma = [\sigma_1 = \{i' = len' = 0, r = 1, r' = 0, \equiv_{i, len}\}, \sigma_2 = \{i' = len', i' \neq 0, \equiv_{r, i, len}\}]$ . As we further advance into the loop,  $\sigma_2$  will maintain equivalence (and in fact will not require widening?) but  $\sigma_1$  will continue to update the part of the state regarding un-tagged variables (since  $r' = 1$  in  $\sigma_1$  and it will not consider any of the commands guarded by  $r'$ ), specifically it will decrement  $i$  continuously, preventing the analysis from reaching fixed point. We would require widening here but using the naive strategy of a complete join will result in aggressive loss of precision, specifically losing all information regarding  $i$ . The problem originates from the fact that prior to widening, we joined sub-states which adhere to two different loop behaviors: one where both P and P2 loop together (that originated from  $len \neq 0$ ) and the other where P2 has existed but P continues to loop ( $len = 0$ ). Ideally, we would like to match these two behaviors and widen them accordingly. We devised a widening strategy that allows to do this as it basically matches sub-states that adhere to the same behavior, or loop-paths. We do this by using guard for the matching. If two sub-states agree on their set of guards, it means they represent the same loop path and can be widened as the latter originated from the former (widening operates

on subsequent iterations). In our example, using this strategy will allow for a precise description of the difference conserved in the state  $\sigma_1 = len = len' = 0, -\inf < i < 0, i' = 0, r = 1, r' = 0$  showing that the original program continued to access the array at illegal  $-\inf < i < 0$  indexes, where P2 stopped and exited. We also define the notion of *Widening Point* meaning at what point in the loop we widen and experiment with applying at back-edges and differencing points. **maybe add a nice picture of behaviors in P and P2 and how they are widened**.

## 5.1 Correlating Abstract State Differencing

Given a state in our correlating domain, we want to determine whether equivalence is kept and if so under which conditions it is kept (for partial equivalence) or determine there is difference and characterize, as much as possible, this difference. As our state may hold several sub-states, each holding different equivalence data, we can provide a much more verbose answer regarding whether equivalence holds. We partition our sub-states according to the set of variables they hold equivalence for and report the state for each equivalence partition class. Since we instrument our correlating program to preserve initial input values, for some of these states we will also be able to report input constraints thus informing the user of the input ranges that maintain equivalence. In the cases where equivalence could not be proved, we report the offending states and apply a differencing algorithm for extraction of the delta. Fig. ?? shows an example of where our analysis is unable to prove equivalence (as it is sound), although part of the state does maintain equivalence (specifically for  $x = 0$ ). This is due to the abstraction being too coarse. We describe an algorithm that given a sub-state  $rd \in RD$ , computes the differentiating part of the sub-state (where correlated variables disagree on values) by splitting it into parts according to equivalence. This is done by treating the relational constraints in our domain as geometrical objects and formulating delta based on that.

**Correlating Abstract State Delta** Given a sub-state  $rd$  and a correspondence  $VC$ , the correlating state delta  $\Delta_A(rd)$ , computes abstract state differentiation over  $rd$ . The result is an abstract state  $\sqsubseteq rd$  approximating all concrete values for variables correlated by  $VC$ , that differ between  $P$  and  $P'$ . Formally, the delta is simply the abstraction of the collecting semantics deltas  $\alpha(\Delta_{CS}^+), \alpha(\Delta_{CS}^-)$  i.e. an abstraction of the set differencing operation. But it is not clear as to how we compute this set differencing on the correlating abstraction. Instead, we consider the geometric representation of the domain and applied operations there for the extraction of delta, as following:

1.  $rd_{\equiv}$  is a state abstracting the concrete states shared by the original and patched program. It is achieved by computing:  $rd_{\equiv} \triangleq rd|_{V=V'} \equiv rd \sqcap \bigwedge \{v = v' | VC(v) = v'\}$ .
2.  $\overline{rd_{\equiv}}$  is the negated state i.e.  $RD \setminus rd_{\equiv}$  and it is computed by negating  $rd_{\equiv}$  (as mentioned before, all logical operations, including negation, are defined on our representation of an abstract state).
3. Eventually:  $\Delta_A(rd) \triangleq rd \sqcap \overline{rd_{\equiv}}$  abstracts all states in  $P \times P'$  that where correlated variables values do not match.
4.  $\Delta_A(rd)^+ = \Delta_A(rd)|_{V'}$  is a projection of the differentiation to display values of  $P'$  alone i.e. "added values".
5.  $\Delta_A(rd)^- = \Delta_A(rd)|_V$  is a projection of the differentiation to display values of  $P$  alone i.e. "removed values".

Applying the algorithm on Fig. ??'s  $P$  and  $P'$  where  $rd = \{retVal' = 2retVal\}$  will result in the following:

1.  $rd_{\equiv} = \{retVal' = 0, retVal = 0\}$ .

Figure 5: Delta computation geometrical representation.

2.  $\overline{rd} = [\{retVal' > 0\}, \{retVal' < 0\}, \{retVal > 0\}, \{retVal < 0\}]$
3.  $\triangle_A(rd) = [\{retVal' = 2retval, retVal' > 0\}, \{retVal' = 2retval, retVal' < 0\}, \{retVal' = 2retval, retVal > 0\}, \{retVal' = 2retval, retVal < 0\}]$
4.  $\triangle_A(rd)^+ = [\{retVal' > 0\}, \{retVal' < 0\}]$
5.  $\triangle_A(rd)^- = [\{retVal > 0\}, \{retVal < 0\}]$

We see that displaying the result in the form of projections is ill-advised as in some states differentiation data is represented by relationships on correlated variables alone, thus projecting will lose all data and we will be left with a less informative result. A geometrical representation of  $\triangle_A$  calculation can be seen in Fig. 5.

---

**Algorithm 1:** Compute Abstract Difference.

---

**Input:** Abstract state  $rd$   
**Output:** Abstract difference -  $\triangle(rd)$   
 $rd_{\equiv} \leftarrow rd \sqcap \bigwedge \{v = v' \mid VC(v) = v'\}$   
 $\overline{rd_{\equiv}} \leftarrow \neg rd_{\equiv} = \{\neg c \mid c \in rd_{\equiv}\}$   
Foreach  $rd_i \in \overline{rd_{\equiv}}$ :  $\triangle(rd) \leftarrow \triangle(rd) \cup (rd \sqcap rd_i)$   
**return**  $\triangle(rd)$

---

From this point forward any mention of 'delta' (denoted  $\triangle$ ) will refer to the correlating abstract state delta (denoted  $\triangle_A$ ). We claim that  $\triangle$  is a correct abstraction for the concrete state delta which allows for a scalable representation of difference we aim to capture.

## 6. Semantics of Correlating Program $P \bowtie P'$

### 6.1 Definition of Difference

### 6.2 General Product Program

A simple approach for a joint analysis is to construct a product program  $P \times P'$  where at every point during the execution we can perform a program step (as defined in Definition ??) of either programs. The product program has a duo-state  $(\sigma, \sigma')$  and each step updates  $(\sigma, \sigma')$  accordingly. The product program can also be seen as a concurrent run  $P || P'$  where every interleaving is possible. The product program emphasizes the fact that, as described in Section 2, the notion of  $\triangle$  is unclear without an established variable and label correspondence. Choosing the location where  $\triangle$  is checked is a key part of identifying differences. Consider Fig. ??, which presents a product automata of the simple program with itself, we see that even in this trivial program, although it is clear that  $\triangle = \emptyset$ , checking for difference in any of the non-correlating states will result in a false difference being reported. As this example demonstrates, selecting a correct label correspondence is crucial for a meaningful delta, we will elaborate on our approach for choosing  $DP$  in ??.

```

1 void foo() {
2   int x = 0;
3 }

```

Figure 6: Program  $P$

### 6.3 Program Correspondence and Differential Points

Selecting the point where  $\triangle$  is computed is vital for precision. As mentioned, a natural selection for diff points would be at the endpoints of traces but that loses meaning under the collecting semantics. A possible translation of this notion under the collecting semantics would be to compute delta between *all* the endpoints of the two programs i.e.  $DP = \{(fin, fin') \mid fin \in exit(P), fin' \in exit(P')\}$  somehow differentiating the final states of the programs. This approach is problematic for two reasons:

1. Comparing all endpoints results in a highly imprecise delta. This is shown by the simple exercise of taking program with 2 endpoints and comparing it with itself.
2. This choice for  $DP$  may result in missing key differences between versions. If at some point during the calculation existed a delta that failed reaching the final state - it will be ignored. An interesting example for this is an array index receiving different bounds after a patch (but later overwritten so that it is not propagated to some final state).

Alternatively, the brute force approach where we might attempt to capture more potential diffs by selecting a diff-point after every line, will result in a highly inaccurate result as, for instance in Fig. ??, many diffs will be reported although there is no difference. Finally, we must be careful with the selection of  $DP$  as it affects the soundness of our analysis: we might miss differences if we did not correctly place diff-points in locations where delta exists. **Our approach employs standard syntactic diff algorithm ?? for producing the correlation. This selection for  $DP$  assures soundness.** The Diff approach works well since two versions of the same software (and especially those that originate from subsequent check-ins to a code repository) are usually similar. Another important factor in the success of the diff is the guarded instruction format for our programs (as defined in Definition ??). Transforming both programs to a our format helps remove a lot of the "noise" that a patch might introduce yet it is superior to low lever intermediate representation as it retains many qualities (such as variable names, conditions, no temporaries, etc.). **See Appendix ?? for examples illustrating said benefits and qualities.** There are alternate ways for creating the correspondence such as **graph equivalence, etc.**, this could be a subject of future research. Calculating delta according to  $DP$  over the product automata is a complex task as it allows both programs to advance independently. We formulate the *correlating program* as a restricted product automata where we advance the programs while keeping the correlation allowing for a superior calculation of delta using our correlating abstract domain later defined in Section 5.

### 6.4 The Correlating Program $P \cup P'$

We will generally describe the process of constructing the correlating program. A more elaborate and formal description of the algorithm can be found in Algorithm ??. The correlating program is an optimized structure where not all pairs of  $(\sigma, \sigma')$  are considered, but only pairs that result from a controlled execution, where correlating instructions (according to  $DP$ ) in  $P$  and  $P'$  will execute together. This will allow for superior precision. As said, the main idea is to create one program which contains both versions. The correlating starts out as (exactly) the older version  $P$  (after being converted to our guarded instruction form). Afterwards a syntactic diff with  $P'$  (also transformed to guarded mode) is computed (the programs are not combined just yet). In fact, this is the point where  $DP$  is created as the diff supplies us with the correlation between labels we desire. Then  $P'$ 's instructions are interleaved into the guarded  $P$  while maintaining the correlation found by the diff (matched instructions will appear consequentially). Just before a

patched instruction is interleaved into the correlating, all variables that appear in it are tagged, as to make sure that the patched instructions will only affect patched variables. Thus we maintain the semantics of running both programs correctly while achieving a new construct that will allow us to analyze change more easily and precisely. **Fig. ?? holds a complete correlating program of the program in Fig. ?? and it's patched version, a graphic description as a controlled automata is shown in Fig. ??**. Note that if we view the general correlating program as a concurrent program, then this optimized program can be viewed as a partial-order reduction applied over the concurrent program. One final observation regarding the correlating program is that it is a legitimate program that can be run to achieve the effect of running both versions. This ability allows us to use dynamic analysis and testing techniques such as fuzzing ?? and directed automated testing ?? which may produce input that lead to states approximated by  $\Delta$ .

### 6.5 Analyzing Correlating Programs

Analysis of a guarded correlating program has certain caveats. In

```

1  1:  guard g = (i>0);
2      if (g) i--;
3      if (g) goto 1;

```

Figure 7: example program illustrating guard analysis caveat

order to correctly analyze the program in Fig. ?? we need our analysis to assume  $(i > 0)$  whenever taking the true branch on the `if (g)` instruction and  $(i \leq 0)$  when taking the false branch. However, since the `i--` instruction invalidates this assumption we would need to update the guard assumption to  $(i > -1)$  which would complicate the analysis as we would need to consider updating the guard assumption while widening etc. Our solution simply incorporates the guard's assumption the first time it encounters the guard and allows it to flow to the rest of the nodes. We are not in danger of losing the assumption during the following join as our join employs a partition-by-equivalence strategy and will not join the two states where  $g, i > 0$  and  $\neg g, i \leq 0$ .

## 7. Evaluation

We mainly tested our tool on the GNU core utilities, differencing versions 6.10 and 6.11. This benchmark included 40 patches where most of the patches (35) were a one-line patch aimed at updating the version information string in the code. Our analysis easily showed equivalence for these programs. About 10 of these patches included actual changes to numerical variables and we were able to precisely describe the difference. We also tested our tool on a few handpicked patches taken from the Linux kernel and the Mozilla Firefox web browser.

We implemented a union compiler named *ucc* which creates union programs from any two C programs as well as a differencing oriented dataflow analysis solver for analyzing union programs, both tools use the LLVM and Clang compiler infrastructure. We analyze C code directly thus benefiting from a low number of variables as there are no temporary values as there might appear in an intermediate representation. We also benefit from our delta being computed over original variables. As mentioned in Section ??, we normalize the input programs before unifying them for a simpler analysis.

Analysis of some of our benchmarks required the use of widening. We applied a basic widening strategy which widens all cfg blocks once reaching a certain threshold. All of our experiments were conducted running on a Intel(R) Core-i7(TM) processor with 4GB.

## 7.1 Results

Table 1: Experimental Results

Name	Domain	#Added	#Deleted	#DiffPoints	#Diffs
dd.i	ppl	52	54	87	4
id.i	ppl	15	6	26	0
pr.i	ppl	10	3	13	1
su.i	ppl	2	2	2	0
env.i	ppl	2	2	3	0
seq.i	ppl	10	10	15	4
nice.i	ppl	3	3	36	0
test.i	ppl	3	3	14	0
chmod.i	ppl	3	3	7	0
nohup.i	ppl	2	2	18	0
paste.i	ppl	24	17	4	0
rmdir.i	ppl	3	0	3	0
users.i	ppl	3	4	30	0
chroot.i	ppl	3	3	22	0
md5sum.i	ppl	15	7	32	28
runcon.i	ppl	2	2	3	0
lbracket.i	ppl	3	3	14	0
setuidgid.i	ppl	7	4	34	12
chown-core.i	ppl	3	3	10	0

Tab. 1 summarizes the results of our analysis. The columns indicate the benchmark name and description, lines of code for the analyzed program, the number of lines added and removed by the patch, the number of diff-points generated, the numerical domain used, and the number of differences found in the analysis (i.e. number of diff-points where  $\Delta \neq \emptyset$ ). In our benchmarks, we focused on computing intra-procedural difference between the two versions of procedures. Procedure calls presented difficulty as they potentially change global variables and local variables through pointers. We overcame this by either (i) assuming equivalence (alone) once we encounter a call to a procedure we already established as equivalent or (ii) warn that all results regarding variables touched by the procedure is un-sound. **In the majority of our benchmarks we identified calls only to library and system procedures thus we could omit their effect as they do not change variables beyond those given as parameter or those being assigned the return value.** All differences reported describe, in constraints over variables, an existing delta at that program point.

**Identifying delta down the line** One advantage of our analysis method is the ability to identify differences in variables that were not directly affected by the patch. Fig. ?? shows part of the `bsd_split_3` function that was patched by the line marked by a comment. Note that although the patch directly restricts the value range of `s_len` to above zero, our analysis is able to identify the effects on the index variable `i` and also report a lost program state of `i > -1` further down the line...

**Non-convex delta**

**Maintaining equivalence in loops**

## 8. Related Work

**Bounded symbolic execution in Clang** As prior work we used the Clang infrastructure [1] static analysis graph reachability engine in order to perform a simple and bounded state differentiation exploration. We used the existing infrastructure and it's abstract representation facilities to simply record every location where the 2 versions of the variables differ. This of course was not sufficient since it only presents a bounded solution and we will show the limitations of this method by example.



```

1 static bool
2 bsd_split_3 (char *s, size_t s_len, unsigned char **hex_digest,
3             size_t i;
4
5     if (s_len == 0) return false; // Patch
6
7     *file_name = s;
8
9     /* Find end of filename. The BSD 'md5' and 'sha1' commands do not
10      * find filenames, so search backwards for the last ')'. */
11
12     i = s_len - 1;
13     while (i && s[i] != ')')
14         i--;
15
16     if (s[i] != ')') return false;
17
18     s[i++] = '\0';
19
20     while (ISWHITE (s[i])) i++;
21
22     if (s[i] != '=') return false;
23     i++;
24
25     while (ISWHITE (s[i])) i++;
26
27     *hex_digest = (unsigned char *) &s[i];
28     return true;
29 }

```

Figure 8: md5sum.c bsd\_split\_3 function patch

**Existing work on patch-based exploit generation** Brumley, Poosankam, Song and Zheng [?] is the prominent work addressing patch-based analysis. We differ from this work in the following aspects:

1. First, the problem definition in said work is different from our own. They aim to find an *exploit* for vulnerabilities fixed by a certain patch. Furthermore, this exploit is defined in relevance to a *security policy* which can differ. While our goals are similar to those of [?], we achieve them by solving 2 extended problems of a) recording the delta between new variable values and old ones and b) producing input from said values. These problems are a superset of the problem described in [?] and solving them has the potential for a much more complete and sound result.
2. We aim to find differentiation between every variable changed by the patch and analyze that differentiation while they concentrate on input sanitation alone. Thus if in the patched program some variable has changed in a way that does not involve input validation, it will be disregarded: for instance if an array index variable  $i$  to a buffer  $B$  is patched by adding an assignment  $i = \text{sizeof}(B) - 1$ , it will be ignored in the previous work while we will record that the old version of  $i$  can no longer have values greater than  $\text{sizeof}(B) - 1$  and may use it for exploit generation.
3. We perform our analysis on the source code of the program and patch instead of the binary. Working on a higher level gives us much more data thus potentially allowing for more results.

Kroening and Heelan [?] main focus was producing an exploit from given input that is known to trigger a bug. No patch is involved in the process. Our goal is to produce said input from the corrected software thus [?] can be used to create an exploit from our results.

Song, Zhang and Sun [?] also relate to the patch-based exploit generation problem but their main focus is on finding similarities between versions of the binary to better couple functions from

the original program with their patched counter-part, a problem that was not addressed in [?]. Also their method of recognizing possible exploits is degenerate and relies on identifying known input validation functions that were added to a certain path - a method that could be easily overcome.

Oh [?] presented a new version for the DarunGrim binary diffing tool aimed at better reviewing patched programs and specifically finding patches with security implications. The goal of the tool is to help researches who manually scan patches for the purpose of producing intrusion prevention system signatures. The tool relies mainly on syntactic analysis of patterns to produce a security implication score for procedures patches making them a candidate for manual inspection. [?] used the DarunGrim binary diffing tool EBDS for their experiment.

Person, Dwyer, Elbaum and Pasareanu [7] introduced an extension and application of symbolic execution techniques that computes a precise behavioral characterization of a program change called differential symbolic execution. As we also implemented bounded symbolic execution as our preliminary work we will discuss this method in comparison to our own.

Godlin and Strichman [?] developed a method for proving the equivalence of similar C programs under certain restrictions based on and existing functional verification tool. This was a basis for future work regarding equivalence and we intend to base our work upon these advances.

Kawaguchi, Lahiri and Rebelo [6] defined the concept of *conditional equivalence* meaning under which conditions (inputs) are 2 different versions of a program equivalent (i.e. produce the same output). Their goal is to keep software changes from breaking procedure contracts and changing module behavior too drastically and they achieve this by computing the conditions under which the behavior is preserved. This work indirectly addresses our problem and we believe we can leverage their techniques for producing the inputs that break the equivalence while focusing on bug triggering rather than contract breaking.

[?]

**Determining corresponding components** As suggested in [?], one possibility is to rely on the editing sequence that creates the new version from the original one. Another option is using various syntactic differencing algorithms as a base for computing correspondence tags.

**their idea for computing correspondence, is to minimize the "size of change". They have two different notions of size of change.**

[?] introduced a correlating heap semantics for verifying linearizability of concurrent programs. In their work, a correlating heap semantics is used to establish correspondence between a concurrent program and a sequential version of the program at specific linearization points.

## A. Appendix

**Correlating Collecting Semantics of a Program Label** Given a paired label  $l_x = (l, l') \in \text{Lab}_x$  in the product program we define the correlating collecting semantics  $CS(l_x) \subseteq \Sigma_{P \times P'}$  as all the correlating states which reach that paired label (i.e. exist in some trace reaching that label), grouped by trace. Formally:

1.  $at_x : (\text{Lab}_x \times \llbracket P \times P' \rrbracket) \rightarrow 2^{\Sigma_{P \times P'}}$  for a given label  $l_x$  and trace  $\pi_x$ , is the set of prefixes of  $\pi_x$  that end in a state labeled  $l_x$  formally:  $at(l_x, \pi_x) \triangleq \{\pi_x^* | \pi_x^* \in \text{pre}_x(\pi_x) \wedge \text{last}_x(\pi_x^*) = l_x\}$ .
2.  $CS(l_x) \triangleq \{\{\text{last}_x(at_x(l_x, \pi_x))\} | \pi_x \in \llbracket P \times P' \rrbracket\}$ .

As mentioned, distinguishing groups of states that arrived from different traces is imperative here, otherwise we could be comparing states that originate from different inputs.

Now we have a collecting semantics that collects dual states from correlating executions only, and distinguishes them by input.

**Collecting Semantics Delta** As discussed in ??, we now calculate difference in locations based on a correspondence of labels  $DP$  in  $P$  and  $P'$ . This matches our collecting semantics as our collecting semantics is defined over paired labels. We compute the collecting semantics delta by applying delta on each  $CS(l_x)$  for each matched pair of label  $(l, l') = l_x \in DP$ . This is similar to the label-based trace delta as now we handle *sets* of correlating states (instead of a series of regular states) and we are encountered with the same difficulty. For example, When we compute the collecting semantics for label  $(lab, lab')$  in the example from Fig. ??, we get  $CS(lab, lab') = \{ \{ \langle x \mapsto 0, i \mapsto 0 \rangle, \langle x' \mapsto 0, i' \mapsto 0 \rangle \}, \{ \langle x \mapsto 1, i \mapsto 0 \rangle, \langle x' \mapsto 1, i' \mapsto 0 \rangle \}, \{ \langle x \mapsto 1, i \mapsto 1 \rangle, \langle x' \mapsto 1, i' \mapsto 1 \rangle \}, \{ \langle x \mapsto 1, i' \mapsto 0 \rangle, \langle x' \mapsto 1, i' \mapsto 1 \rangle \}, \{ \langle x \mapsto 1, i \mapsto 0 \rangle, \langle x' \mapsto 1, i' \mapsto 1 \rangle \}, \{ \langle x \mapsto 1, i \mapsto 1 \rangle, \langle x' \mapsto 1, i' \mapsto 1 \rangle \}, \{ \langle x \mapsto 1, i \mapsto 1 \rangle, \langle x' \mapsto 1, i' \mapsto 2 \rangle \}, \{ \langle x \mapsto 1, i \mapsto 1 \rangle, \langle x' \mapsto 1, i' \mapsto 2 \rangle \}, \dots \}$  which indeed groups together all possible states at the desired label according to originating input, but we are still left with the question of how do we calculate delta on each of these groups (for instance the group originating from input  $x = 1$ ). Again, this decision is determined by the kind of difference we want to capture. As before, we choose to define differentiation according to sets of states: we will divide each group of dou-states into two sets containing  $P$  and  $P'$  states, and try to match each of the  $P$  states with a single  $P'$  state (using  $VC$  of course) and vice versa. Formally:

- for each group of correlating states  $\Sigma_x^i \in CS(l_x)$  (which originated from the same input) we break it into  $\Sigma^i = \{ \sigma = first(\sigma_x) | \sigma_x \in \Sigma_x^i \}$  and  $\Sigma'^i = \{ \sigma' = second(\sigma_x) | \sigma_x \in \Sigma_x^i \}$ .
- as  $\Delta_{CS}$  is directional as well, we denote  $\Delta_{CS}^+(\Sigma_x)$  as states of  $P'$  at label  $l$  that do not exist in  $P$  at label  $P'$  (grouped by originating input) i.e. added states and formally:  $\Delta_{CS}^+ \triangleq \{ \Sigma'^i \setminus \Sigma^i | \Sigma_x^i \in CS(l_x) \}$ . We take notice that the  $\setminus$  operation uses  $VC$  for comparison.
- Similarly:  $\Delta_{CS}^- = \{ \Sigma^i \setminus \Sigma'^i | \Sigma_x^i \in CS(l_x) \}$  for "removed" states.

In the previous example  $CS(lab, lab')$ , the  $x \mapsto 1$  group will be broken into  $\Sigma = \{ \langle x \mapsto 1, i \mapsto 0 \rangle, \langle x \mapsto 1, i \mapsto 1 \rangle \}$  and  $\Sigma' = \{ \langle x' \mapsto 1, i' \mapsto 0 \rangle, \langle x' \mapsto 1, i' \mapsto 1 \rangle, \langle x' \mapsto 1, i' \mapsto 2 \rangle \}$  which will then be compared to produce  $\Delta_{CS}^+[x \mapsto 1] \mapsto \{ \langle x' \mapsto 1, i' \mapsto 2 \rangle \}$  and  $\Delta_{CS}^-[x \mapsto 1] = \emptyset$  which successfully describes the difference here as in  $P'$  a new state is added.

#### A.1 Algorithm 2 : Convert $P$ To Guarded Instruction Format

The algorithm is constructive i.e. it takes a procedure  $P$  and outputs the new lines for a guarded version of  $P$ . The original  $P$  is not part of the output.

- *Stage 0:* Output  $P$ 's signature.
- *Stage 1:* Convert all `while` constructs to `if` and `goto` constructs.
- *Stage 2:* For each non branch instruction  $I$ :  
If  $I$  is a declaration, output it under a new block.  
Otherwise collect all branch conditions  $C$  under which  $I$  executes. Produce the code line: `if (  $\bigwedge C$  )  $I$ ;`

We take a small but sufficient example to demonstrate the algorithm:

```
void foo(char Out[], int n) {
    char In[42];
    int i;
    if (n >= 42) n = 41;
    while (i < n) {
        int j = i + 1;
        In[i] = Out[j];
        i++;
    }
}
```

##### Stage 1 Output:

```
void foo(char Out[], int n) {
    char In[42];
    int i;
    if (n >= 42) n = 41;
1:  if (i < n) {
        int j = i + 1;
        In[i] = Out[j];
        i++;
        goto 1;
    }
}
```

##### Stage 2 Output:

```
void foo(char Out[], int n) {
    char In[42];
    int i;
    if (n >= 42) n = 41;
    {
        int j;
        if (i < n) j = i + 1;
        if (i < n) In[i] = Out[j];
        if (i < n) i++;
        if (i < n) goto 1;
    }
1:  }
```

Figure 9: Algorithm 1 application example

Note that as required, the loop is implemented in means of branches and goto and that the branches do not exceed the 1 level of nesting allowed.

#### A.1.1 Algorithm 2 Assumptions

1. We only deal with `while` loop statements.
2. Logical statement (branch conditions) do not have side-effects i.e. they do not change variable values.

#### A.2 Algorithm 1 Example

We run algorithm 1 on the previous example with the following patch to line 4:

```
- if (n >= 42) n = 41;
+ if (n >= 42) n = 40;
```

We start off from the output of algorithm 2 i.e. a "guarded" program that performs simple buffer handling:

```

void foo(char Out[], int n) {
    char In[42];
    int i;
    if (n >= 42) n = 41;

    {
        int j;
1:      if (i < n) j = i + 1;
        if (i < n) In[i] = Out[j];
        if (i < n) i++;
        if (i < n) goto l;
    }
}

```

The result:

```

void foo(char Out[], int n, char Out'[], int n') {
    char In[42], In'[42];
    int i, i';

    if (n >= 42) n = 41;
    if (n' >= 42) n' = 40;

    {
        int j, j';
1:      if (i < n) j = i + 1;
1':    if (i' < n') j' = j' + 1;
        if (i < n) In[i] = Out[j];
        if (i' < n') In'[i'] = Out'[j'];
        if (i < n) i++;
        if (i' < n') i'++;
        if (i < n) goto l;
        if (i' < n') goto l';
    }
}

```

Figure 10: Algorithm 1 Application Example

## B. Worklist

### B.1 Points to Hammer

1. a special kind of self composition, where correlated steps are kept together. This is particularly important when handling loops.

### B.2 TODO

1. run on uc-klee examples.
2. Consider describing each sub-state as a single (possibly looping) path of execution in both programs that originated from the same input.

### B.3 Questions

1. how come you don't need the "product program"?
2. what are the theorems that you provide? (no reason to have definitions if there are no theorems).
3. what abstract domains can we use as "underlying domains" for our abstraction? Do we have any particular requirements from the abstract domains (one requirement is being relational).
4. what makes a "patched version of a program" different from just saying "a different program"? In other words - what are the requirements on the difference between  $P$  and  $P'$ ?
5. can we claim that our abstraction "forgets" paths along which equivalence is established, but keeps apart paths along which there is a difference, hoping that it will re-converge later?
6. (intuition only) what if we correlate badly and lose soundness? one can propose a 2 "trick" programs that correlating them our way gives a result that loses difference.

7. WHY DO WE CHECK DIFFERENCE ABOVE THE SUB-STATE LEVEL? doesn't that mean we compare different paths? isn't that bad?

## References

- [1] clang: a c language family frontend for llvm, 2007.
- [2] BARTHE, G., D'ARGENIO, P. R., AND REZK, T. Secure information flow by self-composition. In *Proceedings of the 17th IEEE workshop on Computer Security Foundations* (Washington, DC, USA, 2004), CSFW '04, IEEE Computer Society, pp. 100–.
- [3] CHEN, L., MINÉ, A., WANG, J., AND COUSOT, P. Interval polyhedra: An abstract domain to infer interval linear relationships. In *Proceedings of the 16th International Symposium on Static Analysis* (Berlin, Heidelberg, 2009), SAS '09, Springer-Verlag, pp. 309–325.
- [4] GODLIN, B., AND STRICHMAN, O. Regression verification. In *DAC* (2009), pp. 466–471.
- [5] KAWAGUCHI, M., LAHIRI, S. K., AND REBELO, H. Conditional equivalence. Tech. rep., MSR, 2010.
- [6] LAHIRI, S., HAWBLITZEL, C., KAWAGUCHI, M., AND REBÊLO, H. Symdiff: A language-agnostic semantic diff tool for imperative programs. In *CAV* (2012).
- [7] PERSON, S., DWYER, M. B., ELBAUM, S. G., AND PASAREANU, C. S. Differential symbolic execution. In *SIGSOFT FSE* (2008), pp. 226–237.
- [8] RAMOS, D., AND ENGLER, D. Practical, low-effort equivalence verification of real code. In *Computer Aided Verification*, vol. 6806 of *LNCS*. Springer, 2011, pp. 669–685.
- [9] RIVAL, X., AND MAUBORGNE, L. The trace partitioning abstract domain. *ACM Trans. Program. Lang. Syst.* 29, 5 (Aug. 2007).
- [10] TERAUCHI, T., AND AIKEN, A. Secure information flow as a safety problem. In *Proceedings of the 12th international conference on Static Analysis* (Berlin, Heidelberg, 2005), SAS'05, Springer-Verlag, pp. 352–367.